InfluenceIQ: Model & Data Card

SciBlitz AI Challenge 2026 — Team sudo_make_it_work (RUET)

1. Datasets / Data Sources

InfluenceIQ does not use a proprietary team-collected training dataset. It works primarily on public inference-time data gathered during campaign execution and uses that data for scoring and reporting, not for model fine-tuning or retraining.

Source	What it is used for	Access path
Instagram, TikTok, X	Creator discovery and enrichment	Apify actors when configured; public meta-tag fallback otherwise
YouTube	Creator discovery and enrichment	Public channel pages and RSS-based collection
Brave Search API	Campaign-relevant creator discovery	Primary search provider
SerpAPI	Campaign-relevant creator discovery	Fallback search provider
scrape.do / direct HTTP fetch	Article and web-page retrieval	Generic page fetching

2. Pre-trained Models Used

Model	Provider	License	Used for
`mrm8488/bert-tiny-finetuned-sms-spam-detection`	Hugging Face	Apache-2.0	Optional spam / low-quality text backend
`microsoft/deberta-v3-base`	Microsoft / HF	MIT	Optional heavier spam-text adapter available in codebase
`unitary/toxic-bert`	Hugging Face	Apache-2.0	Optional toxicity backend
`roberta-base-openai-detector`	Hugging Face / OpenAI	MIT	Optional AI-generated-text likelihood backend
`llama3.1:8b-instruct`	Meta	Llama 3.1 Community License	Optional Llama-based explanation adapter available in codebase
`openai/gpt-oss-20b:free` or configured OpenRouter model	OpenRouter	Apache-2.0 for listed default model	Query planning / optional LLM tasks
`text-embedding-3-small`	OpenAI via OpenRouter	Proprietary API model	Optional semantic relevance embeddings

All model usage is adapter-based and optional. The default repository path is deterministic-first, and model-backed paths run only when relevant flags, dependencies, and external keys are configured. Some listed models are therefore best understood as available adapters rather than always-on runtime defaults; for example, the heavier DeBERTa spam classifier and the Llama-based explainer both exist in the codebase, but are not implied runtime defaults, and the active LLM backend may instead be selected via environment configuration (such as an OpenRouter-backed model).

3. Known Limitations

- **Language and domain bias:** general-purpose pretrained checkpoints may underperform on non-English, code-mixed, or niche creator content.
- **AIGC-detector uncertainty:** AI-generated-text detection is only a supporting signal, not a reliable standalone verdict.
- **Provider-depth variance:** creator-data quality depends on external provider availability; fallback scraping is shallower than primary provider paths.
- **Inference-only evidence limits:** the system only sees public data it can discover and fetch; private audience quality or hidden business context is outside scope.
- **Optional-model inconsistency:** several integrations exist in the codebase but are not part of the default runtime unless explicitly enabled.

4. Ethical Considerations

- **False positives:** spam, toxicity, or risk flags may incorrectly penalize creators, so the platform should be treated as decision support rather than an automatic exclusion engine.
- **Fairness:** pretrained models may reflect biases from their original training data and may not treat all dialects, communities, or creator styles equally.
- **Transparency:** scores should be interpreted alongside source evidence, provenance, and reasoning instead of as opaque rankings.
- **Privacy boundary:** the system is intended to work on public creator, profile, and content data only; it does not require private messages, passwords, or private audience records.

Repository: github.com/md-tonmoy007/InfluenceIQ